



DAVID R. FRELINGER, KARL P. MUELLER

The AGI Rideout Strategy for Reducing Strategic Risk and Promoting Stability in the Transition to Artificial General Intelligence

For more information on this publication, visit www.rand.org/t/PEA4347-1.

About RAND

RAND is a research organization that develops solutions to public policy challenges to help make communities throughout the world safer and more secure, healthier and more prosperous. RAND is nonprofit, nonpartisan, and committed to the public interest. To learn more about RAND, visit www.rand.org.

Research Integrity

Our mission to help improve policy and decisionmaking through research and analysis is enabled through our core values of quality and objectivity and our unwavering commitment to the highest level of integrity and ethical behavior. To help ensure our research and analysis are rigorous, objective, and nonpartisan, we subject our research publications to a robust and exacting quality-assurance process; avoid both the appearance and reality of financial and other conflicts of interest through staff training, project screening, and a policy of mandatory disclosure; and pursue transparency in our research engagements through our commitment to the open publication of our research findings and recommendations, disclosure of the source of funding of published research, and policies to ensure intellectual independence. For more information, visit www.rand.org/about/research-integrity.

RAND's publications do not necessarily reflect the opinions of its research clients and sponsors.

Published by the RAND Corporation, Santa Monica, Calif.

© 2026 RAND Corporation

RAND® is a registered trademark.

Limited Print and Electronic Distribution Rights

This publication and trademark(s) contained herein are protected by law. This representation of RAND intellectual property is provided for noncommercial use only. Unauthorized posting of this publication online is prohibited; linking directly to its webpage on rand.org is encouraged. Permission is required from RAND to reproduce, or reuse in another form, any of its research products for commercial purposes. For information on reprint and reuse permissions, visit www.rand.org/about/publishing/permissions.

About This Paper

Many of the risks associated with racing toward artificial general intelligence (AGI) are driven or exacerbated by expectations that the first state to achieve it will enjoy an overwhelming and enduring first-mover advantage over its geopolitical rivals. In this paper, drawing inspiration from strategies to stabilize the superpower nuclear rivalry during the Cold War, we describe an AGI Rideout strategy that seeks to make artificial intelligence (AI) primacy a less dominating concern for policymaking in the United States and China. AGI Rideout is not a strategy for winning the race to advanced AI or restricting the pace of progress toward it; Rideout is largely compatible with either of these objectives. Instead, it prioritizes maintaining strategic stability and avoiding major war to maximize decisionmakers' long-term options.

This paper is written for a broad audience of government officials, scholars, and AI theorists and practitioners who are concerned with issues surrounding the future of AI and national security.

Center for the Geopolitics of Artificial General Intelligence

RAND Global and Emerging Risks is a division of RAND that delivers rigorous and objective public policy research on the most consequential challenges to civilization and global security. This work was undertaken by the division's Center for the Geopolitics of Artificial General Intelligence (AGI), which is committed to helping decisionmakers understand, anticipate, and prepare to navigate the national security and geopolitical implications of AGI. The center convenes leading technologists, strategists, economists, political scientists, and outside experts to consider the feasibility and effectiveness of prospective AGI-enabled capabilities; the domestic and international implications of their use; and the strategies and policies that governments, businesses, and civil society could adopt to respond to new realities. For more information, visit www.rand.org/geopolitics-of-agi.

Funding

This effort was independently initiated and conducted within the Center for the Geopolitics of Artificial General Intelligence using income from operations and gifts from RAND supporters, including philanthropic gifts made or recommended by DALHAP Investments Ltd., Ergo Impact, Founders Pledge, Charlottes och Fredriks Stiftelse, Good Ventures, Longview, and Coefficient Giving. RAND donors and grantors have no influence over research findings or recommendations.

Acknowledgments

We are grateful for the valuable contributions to the evolution of this paper made by many of our colleagues, both within and outside RAND, including Lisa Abraham, Jim Baker, Hal Brands, Emma

Borden, Zachary Burdette, Beba Cebralic, Michael Chase, Matt Chessen, Matan Chorev, Richard Danzig, Abby Doll, Kathleen Fisher, Rich Girven, Jacob Heim, Chad Heitzenrater, Dan Hendrycks, Andrew Hoehn, Michael Horowitz, Brian Jackson, Nidhi Kalra, William Kim, Bill Marcellino, Mike Mazarr, Jim Mitre, Joel Predd, Philip Root, and Barry Wilson.

Contents

About This Paper.....	iii
Reducing Strategic Risk and Promoting Stability in the Transition to Artificial General Intelligence.....	1
The Importance of a Plan B in the Race to AGI.....	1
Foundations of a Strategy to Ride Out the AGI Transition	7
Implementing the AGI Rideout Strategy	13
Conclusion.....	18
Abbreviations	19
References.....	20
About the Authors	23

Reducing Strategic Risk and Promoting Stability in the Transition to Artificial General Intelligence

The arrival of artificial general intelligence (AGI) was once a hypothetical and distant prospect; now it is an increasingly urgent policy concern. AGI can be expected to present multiple severe and unfamiliar national security challenges.¹ How should the United States approach the problem of safely navigating the transition to a world of AGI under conditions of deep uncertainty?² In this paper, we propose a strategic approach, AGI Rideout, designed to manage geopolitical risks to the United States in the transition to advanced AI. Rideout is not a strategy for “winning” the technological race to AGI, nor does it preclude seeking to do so. Instead, AGI Rideout is aimed at avoiding the risks of major war or catastrophic damage to vital U.S. security interests that could arise from interstate competition to develop AGI while maximizing the strategic options available to future decisionmakers by protecting U.S. AI development from severe disruption.

The Importance of a Plan B in the Race to AGI

Current U.S. strategy for AI competition rests on a foundational premise that maintaining and extending U.S. technological leadership in AI development will naturally translate into enduring strategic advantage.³ This paradigm assumes that, as AGI replicates and self-improves at machine speed, it will provide transformative economic and military benefits, enabling the state that wins the

¹ Jim Mitre and Joel B. Predd, *Artificial General Intelligence’s Five Hard National Security Problems*, RAND Corporation, PE-A3691-4, February 2025; Zachary Burdette, Dwight Phillips, Jacob L. Heim, Edward Geist, David R. Frelinger, Chad Heitzenrater, and Karl P. Mueller, *How Artificial Intelligence Could Reshape Four Essential Competitions in Future Warfare*, RAND Corporation, RR-A4316-1, 2026. We are using AGI as a shorthand label to refer to advanced artificial intelligence (AI) that equals or exceeds human capabilities at performing many important and complex tasks and is capable of a high degree of self-direction and self-improvement with limited human supervision. AGI could be a complex of many narrower AI systems rather than a unitary AGI, but we expect that achieving this level of performance will depend on AI systems that have significant breadth of application.

² There is a host of uncertainties surrounding AGI, including how soon it will arrive, how it will be developed, what it will look like, and what it will be able to do. The AGI Rideout strategy is based on the assumption that AGI will eventually appear but not on a specific expectation about when this is likely to occur; the fundamental security issues associated with the AGI competition and transition apply whether it arrives in a few years or significantly later. The Rideout perspective leans toward skepticism regarding enthusiastic predictions that AGI will be radically transformative on very short timescales in domains in which human organizations and prerogatives are well entrenched; however, such obstacles to rapid policy change make addressing the challenges and risks associated with AGI a matter of urgency even if the arrival of the technology is not literally imminent. See Richard Danzig, *Artificial Intelligence, Cybersecurity, and National Security: The Fierce Urgency of Now*, RAND Corporation, PE-A4079-1, July 2025.

³ See White House, *Winning the Race: America’s AI Action Plan*, July 2025.

race to AGI to enjoy a decisive first-mover advantage (FMA) over its competitors,⁴ even if it is only modestly ahead of its rivals when it crosses that threshold.⁵

This winner-take-all view of the AI competition is central to policy across virtually every domain of government activity, including industrial policy that promotes domestic AI development, export controls that limit adversary access to advanced semiconductors, aggressive talent recruitment initiatives, and massive federal investment in research and development.⁶ This approach reflects long-standing assumptions about the nature of technological competition in the modern era. For more than a century, the United States has relied on technological superiority as a cornerstone of its national security strategy; this has been reflected in the Cold War–era nuclear arms race, the space race, and competitions to dominate global markets in a host of high-tech industrial export sectors. However, such technological advantages have often proved to be contested, fleeting, or less consequential than expected.⁷ Therefore, the United States should consider and prepare for the possibility that the race to AGI will not go strictly according to plan and thus pursue a strategy for AI competition that hedges against its most serious potential risks.

Fragility of the First-to-AGI Paradigm

The archetypical sprint-to-AGI approach to AI competition calls for national effort to be intensely focused on crossing the AGI threshold as fast as possible and, above all, sooner than China, because of the expectation that doing so will result in a very large FMA. The logic of relying on this approach to protect U.S. interests depends either explicitly or implicitly on at least the following four expectations:

- The U.S. AI enterprise is reliably able (with government support as needed) to deliver and operationalize AGI ahead of rivals' efforts.
- The race toward AGI will not lead to hostile actions that seriously impede U.S. AI development or trigger costly international conflict.⁸

⁴ The AI sector in the United States is not a unified national enterprise over which Washington exerts close and centralized control. The question of the extent to which a U.S.-based firm developing AGI would be equivalent to the United States doing so is open to considerable debate. The discussion in this paper will often refer to AGI as if it were a national asset, but we recognize that, in some instances, this simplification obscures important complexities.

⁵ See, for example, Leopold Aschenbrenner, *Situational Awareness: The Decade Ahead*, June 2024.

⁶ White House, 2025; Michael C. Horowitz and Lauren Kahn, "The Cost of the AGI Delusion," *Foreign Affairs*, September 26, 2025. Prominent advocacy of this approach includes Aschenbrenner, 2024; and Special Competitive Studies Project, "Memos to the President: Artificial General Intelligence (AGI)," 2025. It should be noted that U.S. export control policy has not monotonically adhered to maximizing the U.S. lead over China in AI development. However, see Paul Fraioli, ed., "The US Pivot on Regulating AI Diffusion," *Strategic Comments*, Vol. 31, No. 38, December 2025.

⁷ Hal Brands, *Don't Sweat the AGI Race: Why Concerns About Catastrophic Instability Are Overblown*, RAND Corporation, PE-A4188-1, September 2025.

⁸ Such an absence could result from rivals being unwilling to launch preventive attacks against the United States or such attacks being attempted but turning out to be ineffective and not escalating to major conflict.

- AGI technology will pay off as predicted by its more enthusiastic advocates, quickly increasing the economic and military power of the state that possesses it at a rapid and accelerating pace that will leave even fast-follower rivals far behind.⁹
- The ratio of AGI's benefits to its disruptive economic, social, and other effects will be sufficiently favorable that it is desirable to reach the threshold as quickly as possible.

Even if one considers all of these assumptions to be plausible, a strategy of “getting there first” without a backup plan requires essentially flawless execution across these interconnected elements. If any of them fail to materialize, there is great potential for such an approach to lead to a very undesirable outcome. The United States could fall behind in the race to develop or to adopt and operationalize AGI, then be unready to transition to a rapid pursuit of the new leader or poorly prepared to exploit missteps by a competitor. It could certainly win the race but find that the new technology does not render it safe from its enemies, or it could face conflict against rivals while the race is still underway. Economic and other benefits of AGI could manifest slowly or unevenly, leaving time and space for competitors or imitators to catch up. Finally, a no-holds-barred race-to-AGI approach might ultimately yield benefits for the winner yet drive the geopolitical system or individual nations into an unstable state before these benefits can be enjoyed, motivating extreme actions by adversary or U.S. decisionmakers depending, in part, on their largely untested beliefs about what will happen if or when the other side achieves AGI.

Hedging Against Geopolitical Risks

AGI Rideout focuses on mitigating the geopolitical risks that are connected to the pursuit or achievement of advanced AI. These are not the only potential AGI-related dangers worth addressing (as we discuss in later sections), but they appear, at least in the near term, to have the clearest potential to lead to disastrous outcomes on a large scale for the United States.¹⁰ In broad terms, there are four conspicuous categories of geopolitical risk to address (see Table 1).¹¹

⁹ Depending on one's expectations about the dynamics of an AGI takeoff and diffusion of the technology, sustaining the FMA might or might not require active U.S. government effort.

¹⁰ Among potential AI-enhanced adversaries, China generally looms largest because of its size and technological progress, although it is not the only geopolitical threat that could be magnified by canny use of advanced AI. Although we focus here on AI-enabled risks from China, treating Russia, North Korea, and other state and nonstate threats as lesser-included cases, these other potential adversaries should not be forgotten. In particular, it is worth keeping in mind that Russia's nuclear arsenal (and thus its near-term ability to annihilate the United States) dwarfs that of China.

¹¹ See also Zachary Burdette, Karl P. Mueller, Jim Mitre, and Lily Hoak, “Six Ways AI Could Cause the Next Big War, and Why It Probably Won't,” *Bulletin of the Atomic Scientists*, Vol. 81, No. 4, 2025.

Table 1. Key Potential Geopolitical Risks in the Competition to Develop AGI

Risk Category	Definition	Key Concerns
Catastrophic attack	Adversaries using AI or AGI capabilities to threaten U.S. survival or devastating damage	• Massive cyberattacks • Novel bioweapons • Neutralization of U.S. nuclear deterrent • Most likely during a larger conflict
Preventive attacks	Strikes or other hostile actions to impede U.S. AGI development	• Major attacks on AI labs, data centers, chip foundries, or key personnel
Major regional aggression	AI-enabled conflicts that could lead to war between the United States and a major-power enemy	• Taiwan deterrence failure • Other regional conflicts with potential for escalation
Technological eclipse	Being surpassed in global economic and technological competition	• Undermines U.S. sovereignty or results in catastrophic loss of strategic position

Catastrophic attack risks involve adversaries using AI- or AGI-enabled capabilities to directly threaten U.S. national survival or to inflict physically or psychologically crippling harm against the United States and its population. Concerns about such killing-shot threats tend to center on the use of advanced AI to plan, coordinate, and execute cyberattacks against essential infrastructure and other systems on a massive scale; to develop very powerful yet controllable bioweapons; or to neutralize the U.S. strategic nuclear arsenal through attacks against nuclear command, control, and communications systems or other targets that would leave the United States or its allies open to nuclear blackmail.¹² None of these possibilities appears likely to appeal to Beijing as a bolt-from-the-blue attack option, but an insecure adversary with less to lose and a severe sense of desperation or fear of impending or worsening conflict with the United States might be more inclined to consider such a gambit. There may also be possibilities for using advanced AI to enable more-subtle catastrophic attacks through information warfare or other means.

Preventive attacks to impede U.S. AGI development involve seeking to prevent the United States from achieving a dominant and potentially unassailable position in the international system through possession of AGI (or other advanced AI). These attacks could use either kinetic or nonkinetic weapons to strike data centers and AI labs, chip foundries, critical supporting infrastructure, or key personnel. There are good reasons to expect adversaries to be less enthusiastic about attempting such actions in practice than strategic theorists might anticipate,¹³ but we should not dismiss the possibility that the leaders of a powerful state, if they anticipated that the United States would use an AGI

¹² Mitre and Predd, 2025; James D. Fearon, “One Does Not Simply Dismiss the Nuclear Revolution,” in Jim Mitre, Michael Horowitz, Natalia Henry, Emma Borden, and Joel Predd, eds., *Implications of AGI Race Dynamics on International Security and Stability*, RAND Corporation, PE-A4134-1, 2025.

¹³ Zachary Burdette and Hiwot Demelash, *Evaluating the Risks of Preventive Attack in the Race for Advanced AI*, RAND Corporation, PE-A3691-13, September 2025.

advantage to destroy or subjugate their country, might be willing to run substantial risks to avert such an outcome if they believed they had the means to do so.¹⁴

Major regional aggression may be the most dangerous category of conflict risk associated with advanced AI because of the variety of scenarios that could potentially escalate into costly wars involving the United States and powerful adversaries. Such conflicts could of course occur independently of AI progress. However, new, potent AI-enabled conventional, cyber, or other military capabilities developed or deployed by such states as China or Russia could undermine deterrent stability by making successful aggression against important regional neighbors appear feasible at low expected costs. The resulting risk would extend beyond the local region, including to key U.S. allies and the United States itself. The most obvious such scenario would be AI-facilitated aggression by China against Taiwan (which might begin with actions more limited than a cross-Strait invasion); this would also represent a special case in terms of international economic fallout given the global importance of Taiwan's fragile semiconductor industry. However, the list of plausible regional conflicts with great potential for both escalation and U.S. military involvement could become significantly longer with AI-driven shifts in regional military balances.

Technological eclipse is the final category of geopolitical threat: the possibility of being surpassed and then left behind in a global technological and economic competition to a degree that imperils U.S. control over its own destiny or seriously undermines U.S. sovereignty.¹⁵ Because the United States has economic, military, and technological advantages, this is a prospect that has sometimes appeared remote in discussions of U.S. national security since the end of the Cold War, punctuated by periods of anxiety about the possibility of China's rapid growth bringing the era of U.S. unipolar dominance to an end.¹⁶ However, if AI revolutionizes rates of change in these areas, it could make such relative-decline possibilities more salient, particularly if the arrival of AGI has severe disruptive consequences within the United States to which China might not be equally vulnerable.

Primarily because of the potentially catastrophic consequences of major-power war, the AGI Rideout strategy considers all four of these types of geopolitical risk worth hedging against before, during, and after the arrival of AGI.

Shooting the Rapids, Not Competing in a Regatta

AGI Rideout is compatible with seeking to develop AGI first as long as pursuit of that goal is balanced with measures to hedge against the possibility that one or more of the optimistic accelerationist assumptions fails to pan out. Conversely, Rideout does not reject arguments that the

¹⁴ Karl P. Mueller, *Heeding the Risks of Geopolitical Instability in a Race to Artificial General Intelligence*, RAND Corporation, PE-A3691-12, July 2025a.

¹⁵ This is a higher threshold for geopolitical danger than a vague decline in U.S. international leadership, credibility, or prestige, a prospect that often leads to hand-wringing in the United States. Not having a degree of global preeminence to which past U.S. leaders were accustomed may be uncomfortable, but here the concern is the subset of such conditions under which national security is harmed in fundamental and unambiguous ways.

¹⁶ In contrast, prior to the 1990s, fears of U.S. global power being surpassed by rising rival states or blocs—the Soviet Union, China, the Communist bloc as a whole, even Japan in the 1980s—were far from uncommon. See, for example, A. F. K. Organski, *World Politics*, Knopf, 1960; Paul Kennedy, *The Rise and Fall of the Great Powers*, Random House, 1987; and Ezra F. Vogel, *Japan as Number One: Lessons for America*, Harper & Row, 1985.

potential risks of AGI justify slowing or even halting its pursuit until the safety of the technology can be assured. However, at the most fundamental level, AGI Rideout treats the trajectory of AI progress as a torrent that the U.S. government (or international governance) may be able to influence but will likely be unable to entirely control either in the direction of acceleration or restraint.¹⁷ Unlike some past technological competitions that are often proffered as historical precedents, such as the nuclear weapons revolution or the space race, progress in AI is driven far more by corporations anticipating enormous profits than by demand from governments. Better historical analogies for the race toward AGI—and more troubling ones for advocates of technological restraint—might include the first industrial revolution, European colonialism from the 16th to 18th centuries, the telecommunications revolution, or global climate change. This perspective is reflected in the “Rideout” label: a strategic focus on avoiding being dashed against the rocks as technological logic and market forces sweep the world downstream in the direction of increasingly advanced AI and emerging from the maelstrom with the best feasible set of geopolitical options for the subsequent era.

This is one of several points of contrast between AGI Rideout and the 2025 Mutual Assured AI Malfunction (MAIM) proposal.¹⁸ That concept advocates racing toward more-advanced AI but stopping short of crossing the AGI threshold because of the existential threat that the authors argue the technology would pose if developed, and compelling other states to exercise similar restraint by threatening to destroy their AI development infrastructures should they be reckless enough to insist on trying to develop AGI. The authors also propose increasing the U.S. AI development ecosystem’s vulnerability to preventive attack by rivals, in case this becomes necessary to prevent AGI development in the United States.¹⁹

MAIM is built on multiple questionable assumptions.²⁰ First, it posits that a high degree of political consensus can be formed around the proposition that AGI inherently poses an intolerable existential threat to states that do not have it or to humanity in general. Second, it assumes that major powers will be able to reliably recognize AI development progress that is dangerously close to achieving AGI both at home and abroad so that they can halt their own AI development before crossing that threshold and credibly threaten the use of force to deter their rivals from doing so.²¹ Finally, MAIM argues that, if the threat of such actions does not deter attempts to develop AGI,

¹⁷ Mustafa Suleyman and Michael Bhaskar, *The Coming Wave: Technology, Power, and the 21st Century’s Greatest Dilemma*, Crown, 2023. AGI Rideout does not dismiss the value of U.S. or international governance to manage AI development and diffusion but instead reflects a generally pessimistic perspective about the likelihood of achieving that governance, at least in the near term. Whether prospects for state control will substantially improve over time will presumably depend in part on how the economic and political effects of AI progress affect governmental leverage over the multinational corporations and investors developing the technology. Whether the near term is the only time available to establish a regime to manage the race will largely depend on the technology itself.

¹⁸ Dan Hendrycks, Eric Schmidt, and Alexandr Wang, “Superintelligence Strategy: Expert Version,” *SuperIntelligence—Robotics—Safety & Alignment*, Vol. 2, No. 1, 2025.

¹⁹ Hendrycks, Schmidt, and Wang, 2025.

²⁰ Iskander Rehman, Karl P. Mueller, and Michael J. Mazarr, “Seeking Stability in the Competition for AI Advantage,” commentary, RAND Corporation, March 13, 2025.

²¹ Similar to MAIM, the AGI Rideout strategy depends on having good intelligence about adversary AI development, but this strategy is more modest in its expectations about what can be effectively monitored. MAIM’s assumption that states will be able to recognize the point at which a competitor is on the verge of developing AGI (so that it can be stopped from crossing that threshold at the eleventh hour) strains credulity. As discussed later in this paper, Rideout focuses instead on identifying nascent applications of advanced AI for potential aggression before they can be fielded and used at scale.

limited preventive attacks by other states would be able to stop them without escalating to major war. AGI Rideout takes an opposite approach to minimizing catastrophic risks in the AI race: It considers preventive attacks—which, to be effective, would likely require large-scale missile attacks against high-value infrastructure targets in and around the urban centers of a nuclear-armed superpower—as a threat to mitigate, not a survival strategy to embrace.

Hedging Beyond Geopolitics

Potential conflict between the United States and China—or other adversaries wielding capabilities derived from the AGI research efforts of the technological leaders—is not the only category of threat to hedge against in the transition period leading up to and beyond the appearance of AGI. Instability could also come in the form of economic and social disruption within states. Indeed, it seems certain to do so even if the forms and magnitude of the effects are not yet clear. AGI also poses risks from potential loss-of-control incidents, from minor but perplexing incidents to cataclysmic scenarios.²²

The geopolitical risks of the AGI transition differ from these other types of threats because they involve hostile actions by states, and, even more than the other categories, possibly involve catastrophic outcomes that could materialize well before the actual appearance of AGI. Therefore, AGI Rideout focuses on these perils, which are the responsibility of the U.S. national defense enterprise (broadly defined). However, preparing for and dealing with other aspects of the AGI transition may ultimately prove to be a greater challenge than managing the more familiar geopolitical problem of averting major war between superpowers.

Many of the basic principles of the Rideout strategy approach can also be applied to these other dangers: the need to anticipate hazards before they become severe, to proactively develop countermeasures, and to draw on new technological tools to assist in the effort rather than counting on AGI's machine intelligence to provide solutions to the problems later. Although it is outside the scope of this paper, actions taken to address threats of AGI-driven economic and social disruption, criminal actions enabled by advanced AI, or the consequences of AGI that is misaligned with human welfare can be expected to overlap significantly with actions that might be taken to reduce geopolitical risks.

Foundations of a Strategy to Ride Out the AGI Transition

The basic rationale for the AGI Rideout strategy derives from an assessment that pursuing a first-to-AGI approach to AI competition in isolation is unsatisfactory. This is not to say that being the first state to develop and exploit AGI is not desirable, but relying on winning the sprint to AGI is a fragile, high-risk strategy for assuring U.S. national security interests in this domain. Although the United States may be able to extend its technological lead in AI to arrive at the AGI threshold with a dominant lead in the race and without encountering serious setbacks along the way, this outcome is by

²² Michael J. D. Vermeer, Emily Lathrop, and Alvin Moon, *On the Extinction Risk from Artificial Intelligence*, RAND Corporation, RR-A3034-1, 2025; Eliezer Yudkowsky and Nate Soares, *If Anyone Builds It, Everyone Dies: Why Superhuman AI Would Kill Us All*, Little, Brown, 2025; Benjamin Boudreaux, Michael J. D. Vermeer, Kamaria Horton, and Nidhi Kalra, *The Case for AI Loss of Control Response Planning and an Outline to Get Started*, RAND Corporation, PE-A4232-1, October 2025.

no means assured. Such an approach could go off the rails for a variety of reasons. For example, the U.S. AI industry could fail to continue outperforming Chinese AI developers, perhaps by overconcentrating on a developmental path that proves to be a dead end with respect to developing AGI. Realizing the economic and military benefits of AGI could be a gradual process rather than an explosive one, placing the United States in a position in which a fast-following China could overtake it in either foundational model development or AI adoption. The AI competition could directly or indirectly trigger a costly conflict or preventive action by China or others that seriously impedes U.S. progress on the path to AGI. Finally, the benefits of AGI might be preceded or outweighed by its disruptive economic, social, and other effects to the point that hurrying to reach the threshold as quickly as possible would appear to have been a mistake.

Given these possibilities, the AGI Rideout approach is based on the following propositions that depart significantly from a perspective of first-to-AGI optimism:

- Although U.S. AI innovators have considerable advantages over foreign competitors, making AGI most likely to emerge in the United States first, China or other actors that are following in the race (either by necessity or design) could still find powerful applications for advanced AI that the United States is ill-prepared to deal with.
- The benefits of AGI are very unlikely to manifest themselves overnight in most use cases because of the messy realities of human politics and organizational change, creating opportunities for fast followers to catch up in significant areas of AGI implementation.²³ Among other implications, this suggests that the achievement of AGI will not mark the end of the AI competition.
- It is not certain that achieving AGI first would give the United States (or a rival) an enduring monopoly over the technology, let alone exclusive control that could be maintained without severely restricting deployment of AGI to minimize its diffusion. Fighting major wars to suppress technological challengers might be necessary to defend a monopoly, yet doing so might not be sufficient to stymie challengers' technological progress.
- The potential for AI competition to lead directly or indirectly to serious international instability or conflict escalation is significant enough to take seriously;²⁴ if a war between China and the United States were to occur, its costs would likely outweigh the marginal benefits of early AGI achievement. These risks could result from misperception and miscalculation rather than the true capabilities of AGI and would be heightened by U.S. preventive actions to hamstringing China's AI development infrastructure.
- The state that proves to be most successful in the AGI race may be the one that best manages the disruptive effects of AGI rather than the one that acquires it first.²⁵

²³ Colin H. Kahl and Jim Mitre, "The Real AI Race: America Needs More Than Innovation to Compete with China," *Foreign Affairs*, July 9, 2025.

²⁴ AGI Rideout is not based on an expectation that these risks will be high—China and the United States have powerful incentives not to go to war with each other. Instead, it accepts the premise that even a modest potential for major-power war is sufficient to justify serious effort to mitigate it.

²⁵ Matan Chorev and Joel Predd, "America Should Assume the Worst About AI: How to Plan for a Tech-Driven Geopolitical Crisis," *Foreign Affairs*, July 22, 2025; Michael J. Mazarr, *A New Age of Nations: Power and Advantage in the AI Era*, RAND Corporation, PE-A3691-14, January 2026.

AGI Rideout Strategy Objectives

The AGI Rideout strategy is intended to guide the United States safely through a transition period of indeterminate length encompassing the first wave of pre-AGI and AGI applications in the military domain, as well as new developments potentially enabled by the application of AGI to existing development and production processes. This entails the following three strategic objectives:

- Deter adversary aggression against the United States (catastrophic or preventive attack) by denying expected benefits and raising the direct and indirect costs of disruptive actions directed against U.S. technology development or AI-enabled attacks that are intended to cripple the United States as a geopolitical competitor.
- Counter other geopolitical threats (e.g., major regional aggression) enabled by or associated with adversary diffusion of advanced AI into threatening military capabilities.
- Preserve future option space for U.S. AGI development and deployment.

In the domains of military history and nuclear planning, the need to withstand attacks and modulate responses has often been a cornerstone of strategic stability.²⁶ This principle is equally critical in the context of AGI development and deployment. By preserving the ability to endure and adapt to adversary advances, the United States can mitigate the risks of catastrophic escalation and strengthen strategic stability through deterrence, even if it faces temporary disadvantages in the AGI domain.²⁷

The deterrent elements of the Rideout strategy will be relevant for discouraging attacks on U.S. AI programs by rivals regardless of which AI/AGI development strategy the United States pursues. This approach is also aimed at minimizing any geopolitical FMAs that rivals might anticipate from making progress toward AGI by assuring that the United States will continue to contest in the geopolitical space even if U.S. AGI development programs lag. With respect to the underlying technology competition, the overarching objective of the strategy is keeping the United States in the game. AGI Rideout preserves future U.S. leaders' option space by ensuring that the country retains the capability to develop and deploy AGI effectively as needed to prevent either a catastrophic geopolitical reversal or technological eclipse.

Addressing FMA Beliefs in the AGI Competition

The AGI Rideout strategy prioritizes securing U.S. interests and maximizing prospects for achieving favorable outcomes during the transition to and diffusion of AGI rather than accepting increased risk in pursuit of a decisive victory in the race.²⁸ It is focused on measures to deter AI-enabled aggression by altering adversary expectations that attacking would provide strategic benefits;

²⁶ The July Crisis of 1914 is a classic example of the lack of such flexibility having pathological consequences for strategic stability. See, for example, Stephen Van Evera, *Causes of War: Power and the Roots of Conflict*, Cornell University Press, 1999.

²⁷ This strategy parallels historical military doctrines in which the survivability of forces and the ability to execute controlled responses were central to maintaining a balance of power and preventing unchecked aggression.

²⁸ Although the strategy does not assume that being first to achieve AGI will be either sufficient or necessary for achieving a desirable strategic outcome in the AI competition, it does emphasize the importance of not being left behind by Chinese AI development and deployment to the point of being at a serious strategic disadvantage.

many of these initiatives would also contribute to resilience against the effects of such attacks in the event that deterrence fails. Diminishing adversary expectations or fears of an AGI FMA promises additional benefits both within and beyond the national security sphere.

The prospect of a powerful FMA for the state where AGI is invented—or the state that first implements AGI to reap its potential benefits—casts long shadows over the race toward more-advanced AI. Powerful and troubling implications emerge from expecting that achieving AGI before one's competitors by even a narrow margin will produce explosive economic growth or game-changing military capabilities that will quickly (perhaps permanently) leave rivals trailing far behind as the AGI leader develops, innovates, and proliferates exponentially. Moreover, the dynamics in question come into play even if such expectations ultimately prove to be incorrect:

- AI developers will be incentivized to accelerate their efforts in order to outpace their competitors, even if doing so entails exposing humanity to catastrophic risks as a result of relaxing safety standards or cutting corners with respect to the security of their research and the proliferation of their creations. States may overlook warning signs of instability or system failures in the rush to see new capabilities deployed, resulting in AI systems that are powerful but poorly controlled and fielded before their implications are fully understood.
- States that expect or hope to win the race to AGI will not only be encouraged to accelerate their AI development as much as possible but also to impede the progress of rivals that might outpace them, including through acts of military aggression or other extreme measures that could lead to wider international conflict.
- States that see little hope of being the first to achieve AGI may nevertheless consider taking steps to throw sand in the gears of the race leaders by whatever means are available, including capabilities enabled by less-than-general AI. Alternatively, these states may perceive closing windows of opportunity to settle accounts with their enemies before a first-moving, AGI-enhanced hegemon establishes itself.
- Even short of such dramatic events directly connected to the AI race, expectations of a high-stakes FMA at the AGI threshold can also undermine international stability more broadly by intensifying rivalries and impeding or preventing potentially beneficial cooperation among the leading AI competitors.

The prospect of an FMA does not make preventive attacks inevitable; many factors will encourage restraint, especially regarding committing acts of war.²⁹ However, the larger the FMA appears to be, the faster it is expected to materialize, and the sooner it appears that AGI could arrive, the more powerful these destabilizing forces will be. The AGI Rideout strategy therefore seeks to minimize adversary expectations that achieving AGI would provide a profound FMA, particularly in the military arena. The stability benefits of this do not depend on making AGI appear less than valuable, they depend only on suppressing the appeal of having AGI *first*.

AGI Rideout also has considerable value in minimizing adversary fears regarding the consequences of the United States achieving AGI first, to the extent that such reassurance is possible without undermining the preceding goal. This extent may be quite limited, however, because China

²⁹ Burdette, Mueller, et al., 2025.

and other potential U.S. adversaries are not inclined to view the prospect of a more powerful United States with enthusiasm. Therefore, AGI Rideout treats reassuring Chinese leaders that U.S. development of AGI would not pose a dire threat to China or the Chinese Communist Party as desirable but not essential.³⁰

We have framed AGI Rideout as an approach to managing geopolitical risk during the transition to AGI, but it may contribute to addressing other dangers as well. Dampening FMA concerns between the United States and China could buy time for both states to pursue AGI more safely and better prepare to cope with the societal disruptions that are likely to accompany the arrival and diffusion of AGI. Increased national resilience could also be beneficial if AGI proves difficult to control. Finally, to the extent that FMA concerns impede U.S.-China technological collaboration, ameliorating them should increase the prospects for cooperation between the two states in managing the AGI transition.³¹ Thus, Rideout does not solve the political and technical problems of coordinating safe AGI development, but it should expand the space for those solutions by undermining incentives to race or escalate.

The Nuclear Rideout Analogy and Its Limits

Rideout is a term borrowed from the arena of nuclear strategy. A nuclear rideout doctrine is one that calls for waiting to respond to an enemy nuclear strike until after the attack and its results can be assessed instead of relying on a hair-trigger launch-on-warning or launch-under-attack retaliation policy to deter a nuclear first strike. The AGI Rideout strategy is inspired by this precedent, not based on it, so the analogy is worth unpacking.³²

During the first decade of the Cold War nuclear arms race between the United States and the Soviet Union, the superpowers faced nuclear first-mover and first-strike concerns at several levels. Given a sufficient number of deliverable nuclear weapons and fairly good information about the locations of the adversary's forces, it was reasonable to think that a surprise attack might be able to destroy a large part of the enemy's strategic arsenal, of which the vast majority was deployed on land-based bombers or missiles in unhardened launch sites. This idea created potential temptations to strike first during a crisis and incentivized states that feared such an attack to prepare to respond to an incoming enemy attack by striking back as quickly as possible, before most of the warheads arrived. Reducing losses to a state's nuclear forces by launching on warning would enable a more powerful retaliatory strike. Being prepared to do so would presumably discourage an enemy first strike, but it also increased risks of miscalculation, overreaction, or nuclear use in response to false indications of enemy nuclear attack. Furthermore, forces designed to launch on warning also entailed greater potential for accidental or unauthorized nuclear use. At the strategic level, these offensive-advantage

³⁰ For the AGI Rideout strategy, it is not important (and arguably might not be desirable) for a potential U.S. AGI FMA to be small. It matters more that the prospect of the FMA *appears* small to China (although making appearance depart from reality is difficult in practice). In contrast, AGI Rideout seeks to minimize both the expected and actual magnitude of any AGI FMA that China might achieve.

³¹ Oliver Guest and Kevin Wei, *Bridging the Artificial Intelligence Governance Gap: The United States' and China's Divergent Approaches to Governing General-Purpose Artificial Intelligence*, RAND Corporation, PE-A3703-1, December 2024.

³² For a vastly more detailed treatment of this history, see Lawrence Freedman and Jeffrey Michaels, *The Evolution of Nuclear Strategy*, 4th ed., Palgrave Macmillan, 2019.

dynamics contributed to motivations for nuclear arms racing lest the adversary deploy a nuclear force with sufficient qualitative and quantitative advantages to make a more or less disarming nuclear attack appear feasible.³³

Starting in the late 1950s, however, the strategic landscape was transformed by a set of technological changes that made strategic nuclear weapons less vulnerable to attack, including longer-range bombers that could be based far from the enemy, ballistic missiles in hardened silos (and later on elusive mobile launchers), more-robust nuclear command and control systems, and, especially, the advent of nuclear-powered ballistic missile submarines.³⁴ Although options to strike first or to launch on warning remained, considerable resources and effort were devoted to building U.S. and, later, Soviet nuclear force structures and postures that could allow a strategic attack to be absorbed without the target state being disarmed and a measured response to be executed. Whether that decision time should be used for war termination; assembling a massive, coordinated retaliatory strike by surviving forces; or limited nuclear exchanges in pursuit of desired political or military objective was never settled. However, planners considered this issue, and assets were developed and fielded to give national leaders new options where there were few or none before. In the strategic nuclear competition, these developments weakened fears that either side in the superpower arms race might gain an advantage that it could use it to launch an apocalyptic strike against the other, which, in turn, played a large part in enabling nuclear arms control. Treaty verification did not have to be perfect, it merely had to be good enough to detect cheating on a scale that would substantially undermine first-strike stability. The balance of terror became much less delicate.

The AGI transition that the Rideout strategy seeks to navigate is thus more analogous to the decade-long strategic nuclear transition that made nuclear rideout doctrines possible than it is to the act of riding out an enemy nuclear strike. The objectives of the AGI Rideout strategy resemble those of strategies that sought to stabilize the nuclear standoff between the superpowers by weakening incentives to go to war and prioritizing the pursuit of national survival over military victory,³⁵ but the parallel between the two is metaphorical. We should not look to the Cold War nuclear standoff for policies that can be replicated to stabilize the AGI race, and the fact that the Cold War did not turn hot should not be taken as proof that similar policies can prevent conflict during the transition to AGI. Indeed, it should not even be taken as proof that Cold War nuclear strategy was well-crafted or effective, only that it was not catastrophically unsuccessful. The case for adopting the AGI Rideout strategy lies in the present and future, not the past.

³³ Albert Wohlstetter, "The Delicate Balance of Terror," *Foreign Affairs*, Vol. 37, No. 2, January 1959.

³⁴ This transition can aptly be said to have begun with the analysis of vulnerabilities in U.S. bomber forward basing described in Albert Wohlstetter, Fred Hoffman, R. J. Lutz, and Henry S. Rowen, *Selection and Use of Strategic Air Bases*, RAND Corporation, R-266, 1954. See Fred Kaplan, *The Wizards of Armageddon*, Simon and Schuster, 1983; and Freedman and Michaels, 2019.

³⁵ See Thomas C. Schelling, *Arms and Influence*, rev. ed., Yale University Press, 2008; and Michael Krepon, *Winning and Losing the Nuclear Peace: The Rise, Demise, and Revival of Arms Control*, Stanford University Press, 2021.

Implementing the AGI Rideout Strategy

The AGI Rideout strategy employs a multipronged approach to achieve its objectives and in particular to negate threats posed by adversary AI development. The strategy's key elements are

- **monitoring and adapting.** Anticipate, rapidly assess, and respond to changes or impending changes in adversary force capabilities associated with AI innovation and adoption.
- **offsetting adversary advances.** Counter first-wave AGI-enabled military technologies by rapidly fielding new capabilities when needed.
- **protecting AGI research and development (R&D) infrastructure.** Safeguard critical U.S. AGI research and development infrastructure while transitioning toward more dispersed and more resilient configurations to reduce vulnerability and deter disruptive attacks.
- **completing U.S. AGI development.** Maintain continuity in U.S. AGI research, development, and diffusion by preventing adversary actions against other targets from significantly delaying or disrupting national technological progress.

AGI Rideout is a relatively low-cost strategy so long as AI breakthroughs do not necessitate rapid adjustments based on anticipated Chinese capabilities in the context of incompatible national interests. The key is to implement the strategy in such a manner to allow it to be seen and understood by China, to discourage actions against the United States, and for China to understand that the United States is acting consistently with a posture of restraint toward interference in Chinese AI development. This strategy is not entirely defensive in nature, but it reflects a desire to maintain a basic military balance that favors the United States and to thwart undesirable actions, without threatening previously recognized adversary core interests.

Anticipating and Countering Threats from Adversary AGI Advances

Moving from strategy to execution requires concrete organizational changes, resource allocations, and policy decisions. AGI Rideout requires additions to the existing national security apparatus but not a wholesale replacement of a system that deals with a wide variety of challenges. Although AI developments may ultimately entail major changes to the entire national security enterprise, comprising a vast complex of armed forces, the intelligence community, and supporting contractors and industrial base, the urgency of the situation requires leveraging the existing national security establishment to ensure that the United States can be in position to effectively respond to the emerging challenges.

Central to the AGI Rideout strategy is the principle that adversary pre-AGI and AGI military applications need to be anticipated and effective counters to them need to be developed to neutralize or limit the threats they pose to the United States or other states that are important to it. U.S. leadership in the race to AGI *does not* eliminate this requirement. Such a response offers multiple benefits. Most obviously, it contributes to deterrence by denial. In most cases, it also provides defensive value in the event that deterrence fails or if the capabilities proliferate to actors that are not as susceptible to U.S. deterrence efforts as we would expect China to be. Finally, a general expectation that the military advantages to be gained from employing AGI-enabled capabilities against the United

States and its allies are small should tend to reduce the urgency that China will attach to efforts to field them.

AGI and its precursors could lead to a variety of disruptive military capabilities that have the potential to undermine global or regional strategic stability by making a prospective aggressor think that risking or going to war would be more successful, less dangerous, or less costly than it previously would have been. For example, advanced AI might enable the design, production, and coordinated operation of much larger numbers of sensors and weapons, key improvements to military logistics, more-capable and independent autonomous platforms and weapons, or major improvements in the potential scale and sophistication of cyberattacks. In some cases, such capabilities might be countered by equivalent or superior AGI-based defensive measures, but such symmetry in responding to threats is neither necessary nor sufficient.

To deal with these challenges, modifications to the existing intelligence and response systems in the U.S. national security enterprise are needed. As advancing AI accelerates the potential rate of defense innovation and production, the timeliness and rapid response cycle requirements that the intelligence community and the development and acquisition community face will become increasingly demanding.

We propose establishing two major organizational elements to address these challenges presented by the rapidly changing strategic landscape: a Strategic AI Response Agency (SARA) to address destabilizing AI applications and ensure strategic stability during the transition to advanced AI and AGI, and a National Intelligence Center for AI (NIC-AI) to support intelligence functions. The U.S. Department of Defense Strategic Capabilities Office (SCO) would serve as their interface to the armed services. These specialized entities would complement, not replace, existing structures, while aligning efforts with the AGI Rideout strategy's emphasis on resilience, deterrence, adaptability, and timely production capabilities through development of the defense industrial base (DIB) posture to support this type of activity.

SARA, as proposed here, is designed to counter destabilizing AI applications that threaten to weaken U.S. deterrence and make aggression by China or other adversaries easier or more effective. As existing defense organizations continue their broader missions, SARA would work alongside them with a mission of focusing on deterring AI-enabled adversary aggression through denial, ensuring that adversary objectives are thwarted and the costs of aggression exceed possible gains. Additionally, SARA would establish direct connections to the DIB to ensure the timely production of AI-enabled capabilities and the ability to rapidly scale production as needed to meet emerging threats.

SARA would serve as the Pentagon's focal organization for the development, integration, and coordination of capabilities to address destabilizing AI applications. Analogous to the U.S. Missile Defense Agency, SARA would provide centralized strategic direction; develop and demonstrate AI and AGI capabilities and countermeasures to adversary systems, such as mass autonomous platforms and advanced command, control, communications, computer, cyber and intelligence (C5I) networks; and transition mature solutions to the services and combatant commands. The agency would support deterrence-by-denial strategies that would be implemented by the broader national security community, maintain cognizance and support lifecycle management of AI and AGI systems, and coordinate DIB integration and scaling to meet the demands of this new strategic environment.

Importantly, SARA would focus exclusively on AI- and AGI-related risks and opportunities, enabling faster decisionmaking and operational responses while complementing broader defense initiatives.

NIC-AI would support SARA, serving as a specialized intelligence component focusing on monitoring and probing global AI and AGI developments that could pose risks to international stability. Similar to the role of the Defense Special Missile and Astronautics Center (DEFSMAC)—which monitors foreign missile and space activities through technical intelligence collection and analysis within the broader intelligence community—NIC-AI would operate as a specialized entity leveraging AI-enhanced tradecraft to provide actionable intelligence on specific threats. It would not replace existing intelligence structures but instead complement them with AI-focused capabilities and expertise.

NIC-AI would track adversary AI progress and identify potentially destabilizing applications, using AI tools to process and analyze data from various intelligence sources. The center would share the specialized intelligence it develops with SARA, the SCO, and service-level innovation offices to inform decisionmaking. By leveraging AI tools and maintaining a targeted focus, NIC-AI would provide faster and more-specialized intelligence on AI and AGI-related issues than what can be delivered by such organizations as the Weapons Intelligence, Nonproliferation, and Arms Control Center (WINPAC), which must address numerous other priorities. This efficiency is expected to reduce personnel costs in certain intelligence operations through automation, allowing resources to be focused on higher-order analytical tasks.

The SCO would play a crucial role as an interface between SARA and the military services. Building on its established relationships and expertise in integrating emerging technologies, the SCO would help translate SARA's AI/AGI capabilities into service-specific applications. Rather than creating new bureaucratic structures to perform this function, this approach leverages the SCO's existing connections with service-level innovation offices, such as Army Futures and Concepts Command, AFWERX, and NavalX, to facilitate collaboration and ensure interoperability across domains. The SCO would also ensure that AI and AGI solutions integrate effectively with existing systems, minimizing duplication of effort and accelerating the operationalization of new capabilities.

This approach for countering adversary AGI developments offers several advantages. It provides focused oversight for responding to destabilizing AI applications while preserving existing organizational structures. It enhances intelligence capabilities through NIC-AI's specialized capabilities, improving speed and accuracy in monitoring AI-related threats. The framework achieves certain cost efficiencies through AI-enabled automation of intelligence processes. SARA and NIC-AI's dedicated focus would enable faster responses to AI and AGI developments compared with organizations that have broader mandates, potentially essential for maintaining deterrence-by-denial capabilities in an environment of more—rapidly materializing threat capabilities.

The direct connection between SARA and the DIB would facilitate timely production and scalability of AI-enabled systems, while the SCO's involvement would increase the chances of successful cross-service interoperability of use of the systems through its deep connections to service offices tasked with rapidly fielding new capabilities in their respective military domains. Overall, this approach aligns with AGI Rideout's emphasis on resilience, stability, and adaptability, ensuring proactive countermeasures to emerging threats while operating within the existing national security ecosystem.

Possible impediments to implementing this architecture include bureaucratic resistance to establishing new specialized organizations, such as SARA and NIC-AI; coordination complexities between these entities and existing structures; and resource requirements primarily related to DIB modifications. Additionally, ethical and policy issues that are specific to AI and AGI, particularly those relating to the delegation of decisions regarding the use of force to autonomous systems, may slow certain development and deployment timelines that involve the use of those capabilities. However, this hybrid approach appears well-suited to addressing the unique intelligence and timeliness challenges of responding to emergent AGI capabilities without requiring wholesale retooling of the national security community. Establishing SARA as a specialized agency focused on countering destabilizing AI applications, creating NIC-AI as a complementary intelligence component focusing on emergent AI-based military threats, and leveraging the existing SCO as an interface to the services provides focused oversight for AI and AGI challenges enhances intelligence capabilities through specialized AI-focused collection and analysis. Doing so also aligns with the broader objectives of maintaining stability and deterring aggression in the face of rapidly evolving threats from advanced AI.

Protecting U.S. AGI R&D Infrastructure and Completing AGI Development

In the Rideout strategy, developing defensive measures to increase the resilience of U.S. AGI R&D infrastructure complements the development of active countermeasures to adversary AI advances. This element of the strategy specifically entails protecting nonproliferated infrastructure while transitioning to less vulnerable architectures. This would be done to discourage preventive attacks and other aggressive actions intended to impede U.S. AI development and diffusion by making them appear unlikely to succeed. Measures to safeguard AGI development facilities and essential supporting infrastructure and networks would include dispersal, redundancy, and limited internal hardening to reduce damage in the event of kinetic attacks.³⁶ It is important to recognize that, as with many safety measures, these types of defenses would all come at some cost (either direct expense or costs in lost efficiency), which is why a higher degree of hardening and resilience for AI infrastructure is not already in place. For the same reason, adoption of such initiatives can be expected to require U.S. government involvement.

The same is true for protecting AI R&D personnel against potential attacks and manipulation. Doing so could involve a host of measures, reflecting the wide variety of possible kinetic and, especially, nonkinetic attack vectors by which members of this workforce could be targeted, particularly by an adversary with advanced, AI-enabled capabilities. The initial focus would be on functions that are closely aligned with national security applications and would draw on experience in protecting key personnel who work in sensitive positions in other areas. Lessons from this effort might

³⁶ These defensive measures are intended to both increase the likelihood that protected targets survive or can be rapidly restored and complicate adversary operations in ways that reduce the probability of attack success. By raising costs and forcing attacks to grow in intensity and scale, these measures are designed to make adversary actions more likely to trigger U.S. retaliation or result in unacceptable losses or exposure of attack capabilities, altering the attacker's overall calculus. See Karl P. Mueller, "Averting Attacks Against AGI Development: Three Strategic Approaches," in Jim Mitre, Michael C. Horowitz, Natalia Henry, Emma Borden, and Joel B. Predd, eds., *The Artificial General Intelligence Race and International Security*, PE-A4155-1, September 2025b.

then be applied to other sectors of the AGI research and development ecosystems as particular protective measures prove to be cost-effective and acceptable to the personnel they are designed to protect.

In parallel with these defensive measures, the strategy relies on credible threats of punitive retaliation in response to hostile actions directed against U.S. AI infrastructure and personnel. Such deterrent effects would be provided primarily through traditional military and other instruments of national power, potentially augmented by counters to adversary defensive innovations, as discussed previously. Together, these elements are intended to enable continued progress in U.S. AI development and diffusion with minimal interference from rival states in addition to minimizing a source of potentially ruinous interstate conflict. Prudence dictates that these deterrent approaches, including capabilities for countervailing attacks, be designed to remain viable across a range of futures, including scenarios with and without U.S. AGI leadership.

Managing Threats and Discouraging Escalation

The preceding components of the AGI Rideout strategy are aimed at, among other things, deterring China or other major geopolitical actors from undertaking attacks for the purpose of impeding U.S. progress in the AGI competition. These components do so partly to enable AI development to proceed at whatever pace is deemed to best serve U.S. national interests without undue interference, but also to avoid major wars or more-limited acts of aggression that could escalate to larger armed conflict, given the enormous costs that such actions could entail. Because deterrence involves making aggression appear worse than not attacking, it is most effective when the adversary has some assurance that restraint on its part will lead to a tolerable result.

Achieving this is a matter not only of capabilities and action but also of rhetorical and other signals. As the United States pursues an AGI Rideout strategy, leaders would need to clearly communicate their intentions to potential adversaries regarding U.S. responses to their future behavior and U.S. national goals in the AI competition more broadly.³⁷ This would be done to make clear the resilience of U.S. capabilities (both to assure allies and deter adversaries) and to emphasize the measured and well-reasoned character of U.S. strategies and plans to limit escalation risks. However, when words and actions send different messages, one should expect actions to carry the greater weight, particularly if the recipient is predisposed to distrust the United States because of belief or experience. AGI Rideout therefore focuses on observable capabilities and actions as tools for influencing adversaries' behavior rather than on declarations of intent.

Although Rideout does not emphasize trying to assure competitors and others that the United States would use AGI capabilities benevolently (which is largely beyond the scope of a strategy for the national defense enterprise), as a general rule, such reassurance and reward measures can significantly strengthen deterrence if the recipient views them as credible. Seeking opportunities for mutually beneficial AI-related cooperation between the United States and China in areas of common interest,

³⁷ Note that communicating clearly does not require rigidly or explicitly specifying response thresholds or potential actions. States often make ambiguous threats to leave useful room for maneuver, and a clearly understood threat of uncertain but possible retaliation can often be an effective deterrent.

despite the rivalry between them, could contribute to stability and cost avoidance in the AGI race in addition to its potential direct benefits to the two states.³⁸

A Rideout Strategy for China?

AGI Rideout, as presented here, is a strategy for protecting the security and advancing the interests of the United States (although it would not be the sole beneficiary of adopting such an approach to the AI competition). Similar strategic principles could certainly be embraced by other states as well. Indeed, seeking to ride out the AGI transition may be as much as most middle powers can currently aspire to in a technological race in which the United States and China wield such preponderant weight.

Might an AGI Rideout strategy with Chinese characteristics be attractive in Beijing? China may be doing something along these lines already.³⁹ With worldviews shaped by a long history of technological inferiority to the leading Western powers, albeit one that its leadership seeks to overcome, Chinese leaders appear to be less uncomfortable operating from such a position than their U.S. counterparts typically are. A strategy of maintaining a fast follower position in the race to AGI while preparing to excel at fielding applications of advanced AI is not the same as the AGI Rideout strategy presented here, but it more closely resembles this strategy than an approach to the AI race that is dominated by racing to reach the AI threshold first.

Conclusion

Adopting the AGI Rideout Strategy would represent a fundamental shift in how the United States approaches a technological transformation that could be truly profound and, potentially, very disruptive. Rather than betting everything on winning the race to AGI, Rideout emphasizes building the capabilities needed to compete effectively regardless of who achieves AGI first—and to avert the possibility that the AI competition will lead to disaster prior to reaching that threshold. Recognizing and addressing these risks is strategic prudence in the face of radical uncertainty about how a technology that does not yet exist will transform international security; it does not mean admitting defeat or accepting technological inferiority.

The question is not whether the United States should pursue AGI leadership. It absolutely should, although the optimal trajectory and specific objectives of that pursuit remain open to debate. The more fundamental question is whether the United States will couple that effort with the strategic hedging necessary to ensure U.S. security in a world of radical uncertainty, regardless of how the competition unfolds. The AGI Rideout strategy provides a framework for doing so.

³⁸ Michael S. Chase and William Marcellino, *Incentives for U.S.-China Conflict, Competition, and Cooperation Across Artificial General Intelligence's Five Hard National Security Problems*, RAND Corporation, PE-A4189-1, August 2025.

³⁹ “Xi Jinping’s Plan to Overtake America in AI,” *The Economist*, May 25, 2025. For an alternative argument that lack of visible interest in AGI on the part of the Chinese Communist Party may be deliberate misdirection, see Julian Gewirtz, “Global AI Rivalry Is a Dangerous Game,” *Financial Times*, July 29, 2025.

Abbreviations

AGI	artificial general intelligence
AI	artificial intelligence
DIB	defense industrial base
FMA	first-mover advantage
MAIM	Mutual Assured Artificial Intelligence Malfunction
NIC-AI	National Intelligence Center for AI
R&D	research and development
SARA	Strategic AI Response Agency
SCO	Strategic Capabilities Office

References

- Aschenbrenner, Leopold, *Situational Awareness: The Decade Ahead*, June 2024.
- Boudreaux, Benjamin, Michael J. D. Vermeer, Kamaria Horton, and Nidhi Kalra, *The Case for AI Loss of Control Response Planning and an Outline to Get Started*, RAND Corporation, PE-A4232-1, October 2025. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA4232-1.html>
- Brands, Hal, *Don't Sweat the AGI Race: Why Concerns About Catastrophic Instability Are Overblown*, RAND Corporation, PE-A4188-1, September 2025. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA4188-1.html>
- Burdette, Zachary, and Hiwot Demelash, *Evaluating the Risks of Preventive Attack in the Race for Advanced AI*, RAND Corporation, PE-A3691-13, September 2025. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA3691-13.html>
- Burdette, Zachary, Karl P. Mueller, Jim Mitre, and Lily Hoak, "Six Ways AI Could Cause the Next Big War, and Why It Probably Won't," *Bulletin of the Atomic Scientists*, Vol. 81, No. 4, 2025. As of April 2, 2026:
https://www.rand.org/pubs/external_publications/EP70978.html
- Burdette, Zachary, Dwight Phillips, Jacob L. Heim, Edward Geist, David R. Frelinger, Chad Heitzenrater, and Karl P. Mueller, *How Artificial Intelligence Could Reshape Four Essential Competitions in Future Warfare*, RAND Corporation, RR-A4316-1, 2026. As of March 25, 2026:
https://www.rand.org/pubs/research_reports/RRA4316-1.html
- Chase, Michael S., and William Marcellino, *Incentives for U.S.-China Conflict, Competition, and Cooperation Across Artificial General Intelligence's Five Hard National Security Problems*, RAND Corporation, PE-A4189-1, August 2025. As of August 15, 2025:
<https://www.rand.org/pubs/perspectives/PEA4189-1.html>
- Chorev, Matan, and Joel Predd, "America Should Assume the Worst About AI: How to Plan for a Tech-Driven Geopolitical Crisis," *Foreign Affairs*, July 22, 2025.
- Danzig, Richard, *Artificial Intelligence, Cybersecurity, and National Security: The Fierce Urgency of Now*, RAND Corporation, PE-A4079-1, July 2025. As of February 25, 2026:
<https://www.rand.org/pubs/perspectives/PEA4079-1.html>
- Fearon, James D., "One Does Not Simply Dismiss the Nuclear Revolution," in Jim Mitre, Michael C. Horowitz, Natalia Henry, Emma Borden, and Joel B. Predd, eds., *The Artificial General Intelligence Race and International Security*, PE-A4155-1, September 2025. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA4155-1.html>
- Fraioli, Paul, ed., "The US Pivot on Regulating AI Diffusion," *Strategic Comments*, Vol. 31, No. 38, December 2025.
- Freedman, Lawrence, and Jeffrey Michaels, *The Evolution of Nuclear Strategy*, 4th. ed., Palgrave Macmillan, 2019.

- Gewirtz, Julian, “Global AI Rivalry Is a Dangerous Game,” *Financial Times*, July 29, 2025.
- Guest, Oliver, and Kevin Wei, *Bridging the Artificial Intelligence Governance Gap: The United States’ and China’s Divergent Approaches to Governing General-Purpose Artificial Intelligence*, RAND Corporation, PE-A3703-1, December 2024. As of September 16, 2025:
<https://www.rand.org/pubs/perspectives/PEA3703-1.html>
- Hendrycks, Dan, Eric Schmidt, and Alexandr Wang, “Superintelligence Strategy: Expert Version,” *SuperIntelligence—Robotics—Safety & Alignment*, Vol. 2, No. 1, March 2025.
- Horowitz, Michael C., and Lauren Kahn, “The Cost of the AGI Delusion: By Chasing Superintelligence, America Is Falling Behind in the Real AI Race,” *Foreign Affairs*, September 26, 2025.
- Kahl, Colin H., and Jim Mitre, “The Real AI Race: America Needs More Than Innovation to Compete with China,” *Foreign Affairs*, July 9, 2025.
- Kaplan, Fred, *The Wizards of Armageddon*, Simon and Schuster, 1983.
- Kennedy, Paul, *The Rise and Fall of the Great Powers*, Random House, 1987.
- Krepon, Michael, *Winning and Losing the Nuclear Peace: The Rise, Demise, and Revival of Arms Control*, Stanford University Press, 2021.
- Mazarr, Michael J., *A New Age of Nations: Power and Advantage in the AI Era*, RAND Corporation, PE-A3691-14, January 2026. As of March 7, 2026:
<https://www.rand.org/pubs/perspectives/PEA3691-14.html>
- Mitre, Jim, and Joel B. Predd, *Artificial General Intelligence’s Five Hard National Security Problems*, RAND Corporation, PE-A3691-4, February 2025. As of September 17, 2025:
<https://www.rand.org/pubs/perspectives/PEA3691-4.html>
- Mueller, Karl P., *Heeding the Risks of Geopolitical Instability in a Race to Artificial General Intelligence*, RAND Corporation, PE-A3691-12, July 2025a. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA3691-12.html>
- Mueller, Karl P., “Averting Attacks Against AGI Development: Three Strategic Approaches,” in Jim Mitre, Michael C. Horowitz, Natalia Henry, Emma Borden, and Joel B. Predd, eds., *The Artificial General Intelligence Race and International Security*, PE-A4155-1, September 2025b. As of April 2, 2026:
<https://www.rand.org/pubs/perspectives/PEA4155-1.html>
- Organski, A. F. K., *World Politics*, Knopf, 1960.
- Rehman, Iskander, Karl P. Mueller, and Michael J. Mazarr, “Seeking Stability in the Competition for AI Advantage,” commentary, RAND Corporation, March 13, 2025. As of September 18, 2025:
<https://www.rand.org/pubs/commentary/2025/03/seeking-stability-in-the-competition-for-ai-advantage.html>
- Schelling, Thomas C., *Arms and Influence*, rev. ed., Yale University Press, 2008.
- Special Competitive Studies Project, “Memos to the President: Artificial General Intelligence (AGI),” 2025.
- Suleyman, Mustafa, and Michael Bhaskar, *The Coming Wave: Technology, Power, and the 21st Century’s Greatest Dilemma*, Crown, 2023.
- Van Evera, Stephen, *Causes of War: Power and the Roots of Conflict*, Cornell University Press, 1999.

- Vermeer, Michael J. D., Emily Lathrop, and Alvin Moon, *On the Extinction Risk from Artificial Intelligence*, RAND Corporation, RR-A3034-1, 2025. As of September 18, 2025:
https://www.rand.org/pubs/research_reports/RRA3034-1.html
- Vogel, Ezra F., *Japan as Number One: Lessons for America*, Harper & Row, 1985.
- White House, *Winning the Race: America's AI Action Plan*, July 2025.
- Wohlstetter, Albert, "The Delicate Balance of Terror," *Foreign Affairs*, Vol. 37, No. 2, January 1959.
- Wohlstetter, Albert, Fred Hoffman, R. J. Lutz, and Henry S. Rowen, *Selection and Use of Strategic Air Bases*, RAND Corporation, R-266, 1954. As of February 25, 2026:
<https://www.rand.org/pubs/reports/R0266.html>
- "Xi Jinping's Plan to Overtake America in AI," *The Economist*, May 25, 2025.
- Yudkowsky, Eliezer, and Nate Soares, *If Anyone Builds It, Everyone Dies: Why Superhuman AI Would Kill Us All*, Little, Brown, 2025.

About the Authors

David R. Frelinger is a senior policy researcher at RAND. His work spans an array of issues, including crafting strategies to address the rise of advanced artificial intelligence, countering threats posed by cutting-edge and unconventional weapons, and navigating the complexities of nuclear deterrence and force employment. He has an M.A. in political science.

Karl P. Mueller is a senior political scientist at RAND and a faculty member in the RAND School of Public Policy. His research extends across national security subjects that include deterrence and coercion, airpower theory and history, escalation and crisis stability, nuclear strategy, wargaming, and the strategic implications of advanced artificial intelligence. He holds a Ph.D. in politics.